


SynapseGPT — SFT data flow

All SFT inputs, where they live on Drive, and the tokenize → consolidate → mix → train pipeline.



Which notebook runs which script

Each Colab notebook clones the repo and invokes the scripts below. Run order: data-prep + generators → train.

Notebook (.ipynb)	Scripts it runs (.py)	Purpose / output source
sft_data_prep.ipynb	download_sft_data.py → tokenize_sft_data.py →	Download OTS sources (tulu3, dolly, metamath, opencode, samsum, oasst1, alpaca_gpt4, no_robots), tokenize to ChatML, build the data

	consolidate_sft_data.py	registry.
generate_reasoning_distill_colab.ipynb	generate_reasoning_distill.py	Verified reasoning CoT from DeepSeek V4 Flash (orca + numina, rejection-sampled). → reasoning_distill
generate_tool_use_colab.ipynb	generate_tool_problems.py → generate_tool_use.py (uses tools_runtime.py)	Agentic tool-use traces: python (compute) + search/facts (Wikidata-verified). → tool_use
generate_creative_colab.ipynb	generate_creative.py	Creative writing via teacher + LLM-judge (no gold). Short, expressive, judged. → creative
sft_pre_train.ipynb	sft.py	Train SynapseGPT (~2B) on the weighted SFT_DATA_MIX.
(shared, not run directly)	tools_runtime.py	Sandboxed python + Brave search + tokenizer helpers. Used by generate_tool_use AND the inference tool-loop.